

Total motility

Marie Bisconti, Philippe Grosjean & Elise Hennebert

Introduction

The effect of DMSO and ethanol is evaluated at concentrations from 0 up to 2% on spermatozoa total motility.

DMSO

```
mt_dms0 <- readxl::read_excel("../data/Table S2.xlsx",  
  sheet = "total motility_DMSO")  
mt_dms0$donor <- as.factor(mt_dms0$donor)  
names(mt_dms0) <- c("donor", "conc", "motile", "total")  
mt_dms0$motile_frac <- mt_dms0$motile / mt_dms0$total  
skimr::skim(mt_dms0)
```

Table 1: Data summary

Name	mt_dms0
Number of rows	40
Number of columns	5
Column type frequency:	
factor	1
numeric	4
Group variables	None

Variable type: factor

skim_variable	n_missing	complete_rate	ordered	n_unique	top_counts
donor	0	1	FALSE	8	1: 5, 2: 5, 3: 5, 4: 5

Variable type: numeric

skim_variable	n_missing	complete_rate	mean	sd	p0	p25	p50	p75	p100	hist
conc	0	1	0.72	0.74	0.00	0.10	0.5	1.00	2.00	
motile	0	1	100.65	44.93	41.00	72.00	91.0	119.00	260.00	
total	0	1	131.20	56.45	54.00	92.00	124.5	150.50	301.00	
motile_frac	0	1	0.77	0.09	0.53	0.73	0.8	0.83	0.89	

There are eight donors, no missing data.

```
table(mt_dms0$donor, as.factor(mt_dms0$conc))
```

```

      0 0.1 0.5 1 2
1 1    1    1 1 1
2 1    1    1 1 1
3 1    1    1 1 1
4 1    1    1 1 1
5 1    1    1 1 1
6 1    1    1 1 1
7 1    1    1 1 1
8 1    1    1 1 1

```

The data are balanced with one observation for each concentration and each donor and no missing data.

```
mt_dms0_m1 <- glmer(cbind(motile, total - motile) ~ conc + (conc | donor),
  data = mt_dms0, family = binomial(link = "logit"))
summary(mt_dms0_m1)
```

Generalized linear mixed model fit by maximum likelihood (Laplace
Approximation) [glmerMod]
Family: binomial (logit)
Formula: cbind(motile, total - motile) ~ conc + (conc | donor)

Data: mt_dms0

AIC	BIC	logLik	deviance	df.resid
314.1	322.6	-152.1	304.1	35

Scaled residuals:

Min	1Q	Median	3Q	Max
-3.8305	-0.8161	0.0966	0.7369	5.0822

Random effects:

Groups	Name	Variance	Std.Dev.	Corr
donor	(Intercept)	0.111404	0.3338	
	conc	0.007157	0.0846	0.02

Number of obs: 40, groups: donor, 8

Fixed effects:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	1.39733	0.12833	10.89	< 2e-16 ***
conc	-0.18586	0.05615	-3.31	0.000933 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Correlation of Fixed Effects:

(Intr)
conc -0.220

The model is:


,

$$\begin{aligned} \text{motile}_i &\sim \text{Binomial}(n = 1, \text{prob}_{\text{motile}=1} = \hat{P}) \\ \log \left[\frac{\hat{P}}{1 - \hat{P}} \right] &= \alpha_{j[i]} + \beta_{1j[i]}(\text{conc}) \\ \begin{pmatrix} \alpha_j \\ \beta_{1j} \end{pmatrix} &\sim N \left(\begin{pmatrix} \mu_{\alpha_j} \\ \mu_{\beta_{1j}} \end{pmatrix}, \begin{pmatrix} \sigma_{\alpha_j}^2 & \rho_{\alpha_j \beta_{1j}} \\ \rho_{\beta_{1j} \alpha_j} & \sigma_{\beta_{1j}}^2 \end{pmatrix} \right), \text{ for donor } j = 1, \dots, J \end{aligned} \quad (1)$$

Here is a plot of this model:

```
ggplot(data = mt_dms0) +  
  geom_jitter(aes(x = conc, y = motile_frac * 100, col = donor),  
    width = 0.01) +
```

```
geom_line(aes(x = conc, y = fitted(mt_dms0_m1) * 100, col = donor)) +
labs(x = "[DMSO] (%)", y = "Total motility (%)")
```



Generally, slopes are all negative, suggesting a negative concentration effect. Data are rather widespread. There are definitely shifts in the intercept per donor (different motility at [DMSO] = 0). Slopes seems not too different between donors (can the model be simplified?). Let's check it with a likelihood ratio test:

```
mt_dms0_m2 <- glmer(cbind(motile, total - motile) ~ conc + (1 | donor),
  data = mt_dms0, family = binomial(link = "logit"))
anova(mt_dms0_m1, mt_dms0_m2, refit = FALSE) # Despite the name, it is indeed a LR test
```

Data: mt_dms0

Models:

mt_dms0_m2: cbind(motile, total - motile) ~ conc + (1 | donor)

mt_dms0_m1: cbind(motile, total - motile) ~ conc + (conc | donor)


	npar	AIC	BIC	logLik	deviance	Chisq	Df	Pr(>Chisq)
mt_dms0_m2	3	310.59	315.66	-152.30	304.59			
mt_dms0_m1	5	314.11	322.56	-152.06	304.11	0.48	2	0.7866

The likelihood ratio test does not detect significant differences between the full and simplified models at $\alpha = 5\%$. We could thus use the simplest `mt_dms0_m2` model with only a shift in the slope per donor. This model is:

$$\begin{aligned} \text{motile}_i &\sim \text{Binomial}(n = 1, \text{prob}_{\text{motile}=1} = \hat{P}) \\ \log \left[\frac{\hat{P}}{1 - \hat{P}} \right] &= \alpha_{j[i]} + \beta_1(\text{conc}) \\ \alpha_j &\sim N(\mu_{\alpha_j}, \sigma_{\alpha_j}^2), \text{ for donor } j = 1, \dots, J \end{aligned} \quad (2)$$

Here is a plot of this model:

```
ggplot(data = mt_dms0) +
  geom_jitter(aes(x = conc, y = motile_frac * 100, col = donor),
    width = 0.01) +
  geom_line(aes(x = conc, y = fitted(mt_dms0_m2) * 100, col = donor)) +
  labs(x = "[DMSO] (%)", y = "Total motility (%)")
```



```
summary(mt_dms0_m2)
```

```
Generalized linear mixed model fit by maximum likelihood (Laplace
Approximation) [glmerMod]
Family: binomial ( logit )
Formula: cbind(motile, total - motile) ~ conc + (1 | donor)
Data: mt_dms0
```

AIC	BIC	logLik	deviance	df.resid
310.6	315.7	-152.3	304.6	37

Scaled residuals:

Min	1Q	Median	3Q	Max
-3.7957	-0.8682	0.0641	0.8505	5.0878

Random effects:

Groups Name	Variance	Std.Dev.
donor (Intercept)	0.1166	0.3415

Number of obs: 40, groups: donor, 8

Fixed effects:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	1.40292	0.13047	10.753	<2e-16 ***
conc	-0.19271	0.04519	-4.265	2e-05 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Correlation of Fixed Effects:

	(Intr)
conc	-0.256

The Z test indicates that `conc` is significantly different from zero at $\alpha = 5\%$. However, it is not the best test in the case of a mixed model like here. We prefer to rely on the 95% confidence interval calculated either on the profile, or via parametric bootstrap (and especially the later one):

```
confint(mt_dms0_m2, level = 0.95) # 95% CI based on profile
```

Computing profile confidence intervals ...

	2.5 %	97.5 %
.sig01	0.2143841	0.6264441
(Intercept)	1.1208466	1.6907146
conc	-0.2810892	-0.1038180

```
set.seed(8434)
# 1000x parameter bootstrap
(mt_dmso_m2_conf <- confint(mt_dmso_m2, level = 0.95,
  method = "boot", nsim = 1000L))
```

Computing bootstrap confidence intervals ...

2 warning(s): Model failed to converge: degenerate Hessian with 1 negative eigenvalues (and

	2.5 %	97.5 %
.sig01	0.1220672	0.49749800
(Intercept)	1.1470026	1.66783797
conc	-0.2800799	-0.09954095

Slope for conc is significantly different from zero at $\alpha = 5\%$ because the 95% CI does not contain zero.

Additional verifications

Check if there is not a overdispersion (in this case, a binomial generalized model would not be adequate), $Var(Y) = \varphi Np(1-p)$ with φ , the overdispersion coefficient that has to be close to zero. However, “overdispersion is not estimable (and hence practically irrelevant) for Bernoulli models (= binary data = binomial with $N=1$).”, see [glmmFAQ](#). Thus, it cannot be estimated here.

We could double-check the significance of the slope conc by looking at a likelihood ratio test when dropping conc from the model:

```
#drop1(mt_dmso_m2, scope = "conc")
mt_dmso_m3 <- glmer(cbind(motile, total - motile) ~ 1 + (1 | donor),
  data = mt_dmso, family = binomial(link = "logit"))
anova(mt_dmso_m2, mt_dmso_m3, refit = TRUE)
```

Data: mt_dmso

Models:

mt_dmso_m3: cbind(motile, total - motile) ~ 1 + (1 | donor)

mt_dmso_m2: cbind(motile, total - motile) ~ conc + (1 | donor)

	npar	AIC	BIC	logLik	deviance	Chisq	Df	Pr(>Chisq)
mt_dmso_m3	1	114.70	115.14	-57.35	114.70	0.00	1	1.000
mt_dmso_m2	2	114.70	115.14	-57.35	114.70	0.00	1	1.000

```

mt_dms0_m3      2 326.50 329.88 -161.25   322.50
mt_dms0_m2      3 310.59 315.66 -152.30   304.59 17.906  1   2.32e-05 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```

The model with `conc` is significantly different at α level 5% from a reference model that sets the slope `conc` = 0. There is thus a significant effect of DMSO concentration (confirmation of results obtained from 95% CI).

We also double-check convergence of the model by trying different optimisation engines (just to make sure). First, is there a singularity in the model?

```
isSingular(mt_dms0_m2)
```

```
[1] FALSE
```

... then, a report about the model convergence:

```
mt_dms0_m2_all <- allFit(mt_dms0_m2)
```

```
Loading required namespace: dfoptim
```

```
Loading required namespace: optimx
```

```

bobyqa : [OK]
Nelder_Mead : [OK]
nlminbwrap : [OK]
nmkbw : [OK]
optimx.L-BFGS-B : [OK]
nloptwrap.NLOPT_LN_NELDERMEAD : [OK]
nloptwrap.NLOPT_LN_BOBYQA : [OK]

```

```
summary(mt_dms0_m2_all)
```



```

$which.OK
          bobyqa          Nelder_Mead
          TRUE          TRUE
      nlminbwrap          nmkbw
          TRUE          TRUE
      optimx.L-BFGS-B nloptwrap.NLOPT_LN_NELDERMEAD
          TRUE          TRUE
      nloptwrap.NLOPT_LN_BOBYQA
          TRUE

```

```

$msgs
$msgs$bobyqa
NULL

```

```

$msgs$Nelder_Mead
NULL

```

```

$msgs$nlminbwrap
NULL

```

```

$msgs$nmkbw
NULL

```

```

$msgs$`optimx.L-BFGS-B`
NULL

```

```

$msgs$nloptwrap.NLOPT_LN_NELDERMEAD
NULL

```

```

$msgs$nloptwrap.NLOPT_LN_BOBYQA
NULL

```

```

$fixef
          (Intercept)          conc
bobyqa          1.402918 -0.1927083
Nelder_Mead          1.402924 -0.1927091
nlminbwrap          1.402918 -0.1927084
nmkbw          1.402896 -0.1927420
optimx.L-BFGS-B          1.402907 -0.1927069
nloptwrap.NLOPT_LN_NELDERMEAD          1.402879 -0.1927109
nloptwrap.NLOPT_LN_BOBYQA          1.402918 -0.1927070

```

```

$llik
          bobyqa          Nelder_Mead
        -152.2962        -152.2962
      nlminbwrap          nmkbw
        -152.2962        -152.2962
    optimx.L-BFGS-B nloptwrap.NLOPT_LN_NELDERMEAD
        -152.2962        -152.2962
  nloptwrap.NLOPT_LN_BOBYQA
        -152.2962

```

```

$sdcor
          donor.(Intercept)
bobyqa          0.3415152
Nelder_Mead     0.3415163
nlminbwrap      0.3415153
nmkbw           0.3415363
optimx.L-BFGS-B 0.3415439
nloptwrap.NLOPT_LN_NELDERMEAD 0.3415166
nloptwrap.NLOPT_LN_BOBYQA    0.3415171

```

```

$theta
          donor.(Intercept)
bobyqa          0.3415152
Nelder_Mead     0.3415163
nlminbwrap      0.3415153
nmkbw           0.3415363
optimx.L-BFGS-B 0.3415439
nloptwrap.NLOPT_LN_NELDERMEAD 0.3415166
nloptwrap.NLOPT_LN_BOBYQA    0.3415171

```

```

$times
          user.self sys.self elapsed user.child sys.child
bobyqa          0.049   0.001   0.050         0         0
Nelder_Mead     0.059   0.001   0.060         0         0
nlminbwrap      0.044   0.000   0.045         0         0
nmkbw           0.061   0.001   0.062         0         0
optimx.L-BFGS-B 0.346   0.001   0.346         0         0
nloptwrap.NLOPT_LN_NELDERMEAD 0.055   0.000   0.055         0         0
nloptwrap.NLOPT_LN_BOBYQA    0.040   0.000   0.040         0         0

```

```

$feval
          bobyqa          Nelder_Mead
          74          95

```

```

      nlminbwrap      nmkbw
      NA            104
optimx.L-BFGS-B nloptwrap.NLOPT_LN_NELDERMEAD
      14            88
nloptwrap.NLOPT_LN_BOBYQA
      36

```

```

attr("class")
[1] "summary.allFit"

```


Analysis of the residuals

Let's check how the residuals distribute and if there is homoscedasticity.

```


mt_dms0 <- fortify.merMod(mt_dms0_m2)
ggplot(data = mt_dms0, aes(x = .fitted, y = .screid, col = donor)) +
  geom_point() +
  geom_hline(yintercept = 0) +
  labs(x = "Fitted values", y = "Pearson's residuals")

```



There is one extreme value, but otherwise, residuals seem rather correctly distributed. Linearity is good here.


```
ggplot(data = mt_dmsso, aes(x = .fitted, y = sqrt(abs(.screid)))) +
  geom_point() +
  geom_smooth(method = "loess", formula = y ~ x) +
  geom_hline(yintercept = 0) +
  labs(x = "Fitted values", y = "(|Pearson's residuals|)^.5")
```



Homoscedasticity of the residuals seems here acceptable (the blue curve that is a loess smoothing in the data is relatively horizontal).

Note: according to Gauss-Markov theorem that indicates that linearity, random sample, non-collinearity between predictors, non-correlation between predictors and error term and homoscedasticity are the only requirements for our GLMM, we do not have to check it. Also the model is robust to departure of Normality, and none of the tests we made depend on a Normal distribution of the residuals (we said we don't trust z/t tests and replace them by likelihood ratio tests and parameterized bootstrapped confidence intervals). However, for completeness, here is the quantile-quantile plot of the residuals:

```
ggplot(data = mt_dmsso, aes(sample = .screid)) +
  geom_qq() +
  geom_qq_line() +
  labs(x = "Theoretical quantiles", y = "Pearson's residuals")
```



It appears not too bad, except for our extreme value that is clearly visible here at the top. A Shapiro-Wilk test does not confirm Normality, but we are pretty sure it is caused by the extreme value:

```
shapiro.test(mt_dms0$.screid)
```

Shapiro-Wilk normality test

```
data:  mt_dms0$.screid
W = 0.91868, p-value = 0.007
```

Predictions

The model allows to calculate the drop in total mobility according to DMSO concentration from 0 to 2%. Note that an inverse logit transformation is required. Here is an example:

```
mt_dms0_slope <- c(
  ci95_min  = min(mt_dms0_m2_conf["conc", ]),
  estimate  = fixef(mt_dms0_m2)[["conc"]],
  ci95_max  = max(mt_dms0_m2_conf["conc", ]))
```

```
mt_dmslo_slope
```

```
      ci95_min      estimate      ci95_max  
-0.28007992 -0.19270912 -0.09954095
```

```
#saveRDS(mt_dmslo_slope, "../data/motility_total_DMSO_slope.rds")
```

Let's say we want to calculate the drop in total mobility for various DMSO concentrations between 0 and 2% if the total mobility of a sample without DMSO is 80%. The calculation is:

```
predict_logit <- function(conc, intercept = 1, slopes) {  
  slopes_mat <- matrix(slopes, nrow = 1,  
    dimnames = list(NULL, names(slopes)))  
  data.frame(conc = conc, -intercept +  
    boot::inv.logit(boot::logit(intercept) +  
    conc %*% slopes_mat))  
}  
dmslo_conc <- (0:20) / 10  
mt_dmslo_lost <- predict_logit(dmslo_conc, 0.8, mt_dmslo_slope)  
mt_dmslo_lost
```

	conc	ci95_min	estimate	ci95_max
1	0.0	0.000000000	0.000000000	0.000000000
2	0.1	-0.004518953	-0.003101179	-0.001597412
3	0.2	-0.009113322	-0.006238047	-0.003204342
4	0.3	-0.013783163	-0.009410633	-0.004820795
5	0.4	-0.018528474	-0.012618950	-0.006446775
6	0.5	-0.023349194	-0.015863000	-0.008082285
7	0.6	-0.028245198	-0.019142771	-0.009727328
8	0.7	-0.033216298	-0.022458236	-0.011381905
9	0.8	-0.038262242	-0.025809357	-0.013046015
10	0.9	-0.043382709	-0.029196078	-0.014719660
11	1.0	-0.048577308	-0.032618332	-0.016402837
12	1.1	-0.053845580	-0.036076035	-0.018095543
13	1.2	-0.059186994	-0.039569089	-0.019797777
14	1.3	-0.064600945	-0.043097380	-0.021509532
15	1.4	-0.070086754	-0.046660779	-0.023230804
16	1.5	-0.075643667	-0.050259141	-0.024961586

```

17  1.6 -0.081270853 -0.053892306 -0.026701871
18  1.7 -0.086967406 -0.057560096 -0.028451651
19  1.8 -0.092732341 -0.061262318 -0.030210915
20  1.9 -0.098564595 -0.064998762 -0.031979654
21  2.0 -0.104463027 -0.068769201 -0.033757856

```

```
#saveRDS(mt_dmso_lost, "../data/motility_total_DMSO_lost.rds")
```

This is the lost in total motility that the model predicts.

Ethanol

```

mt_etoh <- readxl::read_excel("../data/Table S2.xlsx",
  sheet = "total motility_EtOH")
mt_etoh$donor <- as.factor(mt_etoh$donor)
names(mt_etoh) <- c("donor", "conc", "motile", "total")
mt_etoh$motile_frac <- mt_etoh$motile / mt_etoh$total
skimr::skim(mt_etoh)

```

Table 4: Data summary

Name	mt_etoh
Number of rows	40
Number of columns	5
Column type frequency:	
factor	1
numeric	4
Group variables	None

Variable type: factor

skim_variable	n_missing	complete_rate	ordered	n_unique	top_counts
donor	0	1	FALSE	8	1: 5, 2: 5, 3: 5, 4: 5

Variable type: numeric

skim_variable	n_missing	complete_rate	mean	sd	p0	p25	p50	p75	p100	hist
conc	0	1	0.72	0.74	0.0	0.10	0.50	1.00	2.00	
motile	0	1	92.67	45.35	39.0	61.50	82.50	104.25	238.00	
total	0	1	127.20	56.93	66.0	87.00	117.00	140.25	301.00	
motile_frac	0	1	0.73	0.12	0.4	0.67	0.78	0.81	0.88	

There are also the same eight donors, no missing data.

```
table(mt_etoh$donor, as.factor(mt_etoh$conc))
```

```

  0 0.1 0.5 1 2
1 1    1    1 1 1
2 1    1    1 1 1
3 1    1    1 1 1
4 1    1    1 1 1
5 1    1    1 1 1
6 1    1    1 1 1
7 1    1    1 1 1
8 1    1    1 1 1

```

The data are balanced with one observation for each concentration and each donor and no missing data.

```
mt_etoh_m1 <- glmer(cbind(motile, total - motile) ~ conc + (conc | donor),
  data = mt_etoh, family = binomial(link = "logit"))
summary(mt_etoh_m1)
```

Generalized linear mixed model fit by maximum likelihood (Laplace

Approximation) [glmerMod]

Family: binomial (logit)

Formula: cbind(motile, total - motile) ~ conc + (conc | donor)

Data: mt_etoh

AIC	BIC	logLik	deviance	df.resid
320.3	328.7	-155.1	310.3	35

Scaled residuals:

Min	1Q	Median	3Q	Max
-----	----	--------	----	-----

-4.9916 -0.7375 0.0028 0.9353 1.8933

Random effects:

Groups	Name	Variance	Std.Dev.	Corr
donor	(Intercept)	0.16977	0.4120	
	conc	0.02106	0.1451	-0.02

Number of obs: 40, groups: donor, 8

Fixed effects:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	1.26479	0.15390	8.219	< 2e-16 ***
conc	-0.29582	0.06886	-4.296	1.74e-05 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Correlation of Fixed Effects:

(Intr)
conc -0.163


The model is:

,

$$\begin{aligned} \text{motile}_i &\sim \text{Binomial}(n = 1, \text{prob}_{\text{motile}=1} = \hat{P}) \\ \log \left[\frac{\hat{P}}{1 - \hat{P}} \right] &= \alpha_{j[i]} + \beta_{1j[i]}(\text{conc}) \\ \begin{pmatrix} \alpha_j \\ \beta_{1j} \end{pmatrix} &\sim N \left(\begin{pmatrix} \mu_{\alpha_j} \\ \mu_{\beta_{1j}} \end{pmatrix}, \begin{pmatrix} \sigma_{\alpha_j}^2 & \rho_{\alpha_j \beta_{1j}} \\ \rho_{\beta_{1j} \alpha_j} & \sigma_{\beta_{1j}}^2 \end{pmatrix} \right), \text{ for donor } j = 1, \dots, J \end{aligned} \quad (3)$$

Here is a plot of this model:

```
ggplot(data = mt_eto) +
  geom_jitter(aes(x = conc, y = motile_frac * 100, col = donor),
    width = 0.01) +
  geom_line(aes(x = conc, y = fitted(mt_eto_m1) * 100, col = donor)) +
  labs(x = "[ethanol] (%)", y = "Total motility (%)")
```



Generally, slopes are all negative, suggesting a negative concentration effect. Data are rather widespread. There are definitely shifts in the intercept per donor (different motility at [ethanol] = 0). Slopes seems more different between donors than for DMSO. However, we will also check if the model can be simplified using a likelihood ratio test:

```
mt_etoh_m2 <- glmer(cbind(motile, total - motile) ~ conc + (1 | donor),
  data = mt_etoh, family = binomial(link = "logit"))
anova(mt_etoh_m1, mt_etoh_m2, refit = FALSE) # Despite the name, it is indeed a LR test
```

Data: mt_etoh

Models:

```
mt_etoh_m2: cbind(motile, total - motile) ~ conc + (1 | donor)
```

```
mt_etoh_m1: cbind(motile, total - motile) ~ conc + (conc | donor)
```

	npar	AIC	BIC	logLik	deviance	Chisq	Df	Pr(>Chisq)
mt_etoh_m2	3	319.58	324.65	-156.79	313.58			
mt_etoh_m1	5	320.28	328.72	-155.14	310.28	3.302	2	0.1919


The likelihood ratio test does not detect significant differences between the full and simplified models at $\alpha = 5\%$. We could thus use the simplest `mt_etoh_m2` model with only a shift in the slope per donor. This model is:

,

$$\begin{aligned}
\text{motile}_i &\sim \text{Binomial}(n = 1, \text{prob}_{\text{motile}=1} = \hat{P}) \\
, \log \left[\frac{\hat{P}}{1 - \hat{P}} \right] &= \alpha_{j[i]} + \beta_1(\text{conc}) \\
\alpha_j &\sim N(\mu_{\alpha_j}, \sigma_{\alpha_j}^2), \text{ for donor } j = 1, \dots, J
\end{aligned} \tag{4}$$

Here is a plot of this model:

```
ggplot(data = mt_etoh) +
  geom_jitter(aes(x = conc, y = motile_frac * 100, col = donor),
    width = 0.01) +
  geom_line(aes(x = conc, y = fitted(mt_etoh_m2) * 100, col = donor)) +
  labs(x = "[etoh] (%)", y = "Total motility (%)")
```



```
summary(mt_etoh_m2)
```

Generalized linear mixed model fit by maximum likelihood (Laplace
Approximation) [glmerMod]
Family: binomial (logit)
Formula: cbind(motile, total - motile) ~ conc + (1 | donor)

Data: mt_eto

AIC	BIC	logLik	deviance	df.resid
319.6	324.6	-156.8	313.6	37

Scaled residuals:

Min	1Q	Median	3Q	Max
-5.0521	-0.5711	0.0252	0.9661	2.8884

Random effects:

Groups	Name	Variance	Std.Dev.
donor	(Intercept)	0.1794	0.4236

Number of obs: 40, groups: donor, 8

Fixed effects:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	1.26983	0.15723	8.076	6.69e-16 ***
conc	-0.30297	0.04266	-7.103	1.22e-12 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Correlation of Fixed Effects:

	(Intr)
conc	-0.209

The Z test indicates that conc is significantly different from zero at $\alpha = 5\%$. However, it is not the best test in the case of a mixed model like here. We prefer to rely on the 95% confidence interval calculated either on the profile, or via parametric bootstrap (and especially the later one):

```
confint(mt_eto_m2, level = 0.95) # 95% CI based on profile
```

Computing profile confidence intervals ...

	2.5 %	97.5 %
.sig01	0.2695168	0.7719890
(Intercept)	0.9261419	1.6176274
conc	-0.3865530	-0.2192498

```

set.seed(535)
# 1000x parameter bootstrap
(mt_etoh_m2_conf <- confint(mt_etoh_m2, level = 0.95,
  method = "boot", nsim = 1000L))

```

Computing bootstrap confidence intervals ...

4 warning(s): Model failed to converge with max|grad| = 0.00611004 (tol = 0.002, component 1)

	2.5 %	97.5 %
.sig01	0.1811891	0.6041279
(Intercept)	0.9728800	1.5769631
conc	-0.3882783	-0.2237793

We had one model with singularity among the 1000, not a big problem (we may ignore this warning). Slope for `conc` is significantly different from zero at $\alpha = 5\%$ because the 95% CI does not contain zero.

Additional verifications

We could double-check the significance of the slope `conc` by looking at a likelihood ratio test when dropping `conc` from the model:

```

#drop1(mt_etoh_m2, scope = "conc")
mt_etoh_m3 <- glmer(cbind(motile, total - motile) ~ 1 + (1 | donor),
  data = mt_etoh, family = binomial(link = "logit"))
anova(mt_etoh_m2, mt_etoh_m3, refit = FALSE)

```

Data: mt_etoh

Models:

mt_etoh_m3: cbind(motile, total - motile) ~ 1 + (1 | donor)

mt_etoh_m2: cbind(motile, total - motile) ~ conc + (1 | donor)

	npar	AIC	BIC	logLik	deviance	Chisq	Df	Pr(>Chisq)
mt_etoh_m3	2	367.46	370.84	-181.73	363.46			
mt_etoh_m2	3	319.58	324.65	-156.79	313.58	49.88	1	1.635e-12 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

The model with `conc` is significantly different at α level 5% from a reference model that sets the slope `conc` = 0. There is thus a significant effect of ethanol concentration (confirmation of results obtained from 95% CI).

We also double-check convergence of the model by trying different optimisation engines (just to make sure). First, is there a singularity in the model?

```
isSingular(mt_etoh_m2)
```

```
[1] FALSE
```

... then, a report about the model convergence:

```
mt_etoh_m2_all <- allFit(mt_etoh_m2)
```

```
bobyqa : [OK]
Nelder_Mead : [OK]
nlminbwrap : [OK]
nmkbw : [OK]
optimx.L-BFGS-B : [OK]
nloptwrap.NLOPT_LN_NELDERMEAD : [OK]
nloptwrap.NLOPT_LN_BOBYQA : [OK]
```

```
summary(mt_etoh_m2_all)
```

```
$which.OK
              bobyqa              Nelder_Mead
              TRUE              TRUE
nlminbwrap              nmkbw
              TRUE              TRUE
optimx.L-BFGS-B nloptwrap.NLOPT_LN_NELDERMEAD
              TRUE              TRUE
nloptwrap.NLOPT_LN_BOBYQA
              TRUE

$msgs
$msgs$bobyqa
NULL
```

\$msgs\$Nelder_Mead
NULL

\$msgs\$nlminbwrap
NULL

\$msgs\$nmkbw
NULL

\$msgs\$`optimx.L-BFGS-B`
NULL

\$msgs\$nlptwrap.NLOPT_LN_NELDERMEAD
NULL

\$msgs\$nlptwrap.NLOPT_LN_BOBYQA
NULL

\$fixef	(Intercept)	conc
bobyqa	1.269828	-0.3029685
Nelder_Mead	1.269826	-0.3029668
nlminbwrap	1.269829	-0.3029702
nmkbw	1.269732	-0.3029680
optimx.L-BFGS-B	1.269829	-0.3029684
nlptwrap.NLOPT_LN_NELDERMEAD	1.269756	-0.3029708
nlptwrap.NLOPT_LN_BOBYQA	1.269828	-0.3029680

\$llik	bobyqa	Nelder_Mead
	-156.7893	-156.7893
	nlminbwrap	nmkbw
	-156.7893	-156.7893
	optimx.L-BFGS-B	nlptwrap.NLOPT_LN_NELDERMEAD
	-156.7893	-156.7893
	nlptwrap.NLOPT_LN_BOBYQA	
	-156.7893	

\$sdcor	donor.(Intercept)
bobyqa	0.4235709
Nelder_Mead	0.4235679

```

nlminbwrap                0.4235741
nmkbw                     0.4235690
optimx.L-BFGS-B           0.4235601
nloptwrap.NLOPT_LN_NELDERMEAD 0.4235924
nloptwrap.NLOPT_LN_BOBYQA  0.4235698

$theta
                                donor.(Intercept)
bobyqa                        0.4235709
Nelder_Mead                   0.4235679
nlminbwrap                    0.4235741
nmkbw                         0.4235690
optimx.L-BFGS-B              0.4235601
nloptwrap.NLOPT_LN_NELDERMEAD 0.4235924
nloptwrap.NLOPT_LN_BOBYQA    0.4235698

$times
                                user.self sys.self elapsed user.child sys.child
bobyqa                        0.044      0    0.044      0      0
Nelder_Mead                   0.056      0    0.056      0      0
nlminbwrap                    0.049      0    0.049      0      0
nmkbw                         0.057      0    0.057      0      0
optimx.L-BFGS-B              0.344      0    0.345      0      0
nloptwrap.NLOPT_LN_NELDERMEAD 0.050      0    0.050      0      0
nloptwrap.NLOPT_LN_BOBYQA    0.039      0    0.039      0      0

$feval
                                bobyqa                Nelder_Mead
                                65                    88
                                nlminbwrap              nmkbw
                                NA                      100
                                optimx.L-BFGS-B nloptwrap.NLOPT_LN_NELDERMEAD
                                14                    92
                                nloptwrap.NLOPT_LN_BOBYQA
                                31


attr(,"class")
[1] "summary.allFit"

```

Analysis of the residuals


Let's check how the residuals distribute and if there is homoscedasticity.


```
mt_etoh <- fortify.merMod(mt_etoh_m2)
ggplot(data = mt_etoh, aes(x = .fitted, y = .scre resid, col = donor)) +
  geom_point() +
  geom_hline(yintercept = 0) +
  labs(x = "Fitted values", y = "Pearson's residuals")
```



There is one extreme value (less extreme than for DMSO), but otherwise, residuals seem rather correctly distributed. Linearity is good here.


```
ggplot(data = mt_etoh, aes(x = .fitted, y = sqrt(abs(.scre resid)))) +
  geom_point() +
  geom_smooth(method = "loess", formula = y ~ x) +
  geom_hline(yintercept = 0) +
  labs(x = "Fitted values", y = "(|Pearson's residuals|)^.5")
```



Homoscedasticity of the residuals seems here acceptable (the blue curve that is a loess smoothing in the data is relatively horizontal).

With the same remark as for DMSO, here is the quantile-quantile plot of the residuals:

```
ggplot(data = mt_etoh, aes(sample = .screid)) +
  geom_qq() +
  geom_qq_line() +
  labs(x = "Theoretical quantiles", y = "Pearson's residuals")
```



It appears not too good. A Shapiro-Wilk test indicates mild non Normality (with caution because this test tends to be conservative):

```
shapiro.test(mt_etoh$.scre resid)
```

Shapiro-Wilk normality test

```
data:  mt_etoh$.scre resid
W = 0.94028, p-value = 0.03534
```

Predictions

The model allows to calculate the drop in total mobility according to ethanol concentration from 0 to 2%. Note that an inverse logit transformation is required. Here is an example:

```
mt_etoh_slope <- c(
  ci95_min  = min(mt_etoh_m2_conf["conc", ]),
  estimate  = fixef(mt_etoh_m2)[["conc"]],
  ci95_max  = max(mt_etoh_m2_conf["conc", ]))
mt_etoh_slope
```

```

      ci95_min   estimate   ci95_max
-0.3882783 -0.3029708 -0.2237793

```

```

#saveRDS(mt_etoh_slope, "../data/motility_total_ETOH_slope.rds")

```

Let's say we want to calculate the drop in total mobility for various ethanol concentrations between 0 and 2% if the total mobility in a sample without ethanol is 80%. The calculation is:

```

predict_logit <- function(conc, intercept = 1, slopes) {
  slopes_mat <- matrix(slopes, nrow = 1,
    dimnames = list(NULL, names(slopes)))
  data.frame(conc = conc, -intercept +
    boot::inv.logit(boot::logit(intercept) +
      conc %*% slopes_mat))
}
etoh_conc <- (0:20) / 10
mt_etoh_lost <- predict_logit(etoh_conc, 0.8, mt_etoh_slope)
mt_etoh_lost

```

	conc	ci95_min	estimate	ci95_max
1	0.0	0.000000000	0.000000000	0.000000000
2	0.1	-0.006284872	-0.004891619	-0.003604517
3	0.2	-0.012714728	-0.009871491	-0.007257167
4	0.3	-0.019289630	-0.014939680	-0.010957988
5	0.4	-0.026009417	-0.020096167	-0.014706993
6	0.5	-0.032873700	-0.025340853	-0.018504173
7	0.6	-0.039881848	-0.030673549	-0.022349493
8	0.7	-0.047032985	-0.036093983	-0.026242894
9	0.8	-0.054325977	-0.041601787	-0.030184290
10	0.9	-0.061759427	-0.047196504	-0.034173569
11	1.0	-0.069331667	-0.052877577	-0.038210591
12	1.1	-0.077040754	-0.058644355	-0.042295190
13	1.2	-0.084884466	-0.064496085	-0.046427171
14	1.3	-0.092860294	-0.070431914	-0.050606312
15	1.4	-0.100965441	-0.076450885	-0.054832359
16	1.5	-0.109196822	-0.082551937	-0.059105030
17	1.6	-0.117551062	-0.088733904	-0.063424016
18	1.7	-0.126024496	-0.094995513	-0.067788973
19	1.8	-0.134613171	-0.101335385	-0.072199529
20	1.9	-0.143312852	-0.107752031	-0.076655281

```
21 2.0 -0.152119022 -0.114243858 -0.081155794
```

```
#saveRDS(mt_etoh_lost, "../data/motility_total_ETOH_lost.rds")
```

This is the lost in total mobility that the model predicts.

General informations

```
sessionInfo()
```

R version 4.1.3 (2022-03-10)

Platform: x86_64-apple-darwin17.0 (64-bit)

Running under: macOS Big Sur/Monterey 10.16

Matrix products: default

LAPACK: /Library/Frameworks/R.framework/Versions/4.1/Resources/lib/libRlapack.dylib

locale:

```
[1] en_US.UTF-8/en_US.UTF-8/en_US.UTF-8/C/en_US.UTF-8/en_US.UTF-8
```

attached base packages:

```
[1] stats      graphics  grDevices  utils      datasets  methods    base
```

other attached packages:

```
[1] ggplot2_3.3.5 lme4_1.1-29  Matrix_1.4-1
```

loaded via a namespace (and not attached):

```
[1] tidyr_1.2.0      jsonlite_1.8.0    splines_4.1.3
[4] equationomatic_0.3.1 shiny_1.7.1        assertthat_0.2.1
[7] highr_0.9         broom.mixed_0.2.9.4 cellranger_1.1.0
[10] yaml_2.3.5        globals_0.14.0     numDeriv_2016.8-1.1
[13] pillar_1.7.0      backports_1.4.1    lattice_0.20-45
[16] glue_1.6.2        digest_0.6.29      promises_1.2.0.1
[19] minqa_1.2.4       colorspace_2.0-3   dfoptim_2020.10-1
[22] htmltools_0.5.2   httpuv_1.6.5       pkgconfig_2.0.3
[25] broom_0.8.0       listenv_0.8.0      purrr_0.3.4
[28] xtable_1.8-4      scales_1.2.0        later_1.3.0
[31] tibble_3.1.6      mgcv_1.8-40         generics_0.1.2
[34] farver_2.1.0      ellipsis_0.3.2     withr_2.5.0
[37] furrr_0.2.3       repr_1.1.4          skimr_2.1.4
```

[40]	cli_3.2.0	magrittr_2.0.3	crayon_1.5.1
[43]	readxl_1.4.0	mime_0.12	evaluate_0.15
[46]	fansi_1.0.3	future_1.24.0	parallelly_1.31.0
[49]	nlme_3.1-157	MASS_7.3-56	forcats_0.5.1
[52]	tools_4.1.3	lifecycle_1.0.1	stringr_1.4.0
[55]	munsell_0.5.0	compiler_4.1.3	rlang_1.0.2
[58]	grid_4.1.3	nloptr_2.0.0	rstudioapi_0.13
[61]	base64enc_0.1-3	labeling_0.4.2	rmarkdown_2.13
[64]	boot_1.3-28	gtable_0.3.0	codetools_0.2-18
[67]	DBI_1.1.2	R6_2.5.1	knitr_1.38
[70]	dplyr_1.0.8	optimx_2021-10.12	fastmap_1.1.0
[73]	utf8_1.2.2	stringi_1.7.6	parallel_4.1.3
[76]	Rcpp_1.0.8.3	vctrs_0.4.1	tidyselect_1.1.2
[79]	xfun_0.30		